

Radio Galaxy Zoo: Machine learning methods for radio source host galaxy cross-identification

M. J. Alger¹, J. K. Banfield^{1,2}, C. S. Ong^{4,5}, O. I. Wong^{2,6}, L. Rudnick⁷, R. P. Norris^{3,8}, others

¹Research School of Astronomy and Astrophysics, The Australian National University, Canberra, ACT 2611, Australia

²ARC Centre of Excellence for All-Sky Astrophysics (CAASTRO)

³Western Sydney University, Locked Bag 1797, Penrith South, NSW 1797, Australia

⁴Data61, CSIRO, 7 London Circuit, Canberra ACT 2601, Australia

⁵Research School of Computer Science, The Australian National University, Canberra, ACT 2601, Australia

⁶International Centre for Radio Astronomy Research-M468, The University of Western Australia, 35 Stirling Hwy, Crawley, WA 6009, Australia

⁷Minnesota Institute for Astrophysics, University of Minnesota, 116 Church St. SE, Minneapolis, MN 55455

⁸CSIRO Astronomy & Space Science, PO Box 76, Epping, NSW 1710, Australia

Accepted XXX. Received XXX

ABSTRACT

We present a machine learning approach for the task of determining the infrared host galaxies of radio emissions detected in wide-area radio surveys. We test our methods on the radio sources detected in the 1.4 GHz Australian Telescope Large Area Survey (ATLAS) observations of the *Chandra* Deep Field - South (CDFs) and the ESO Large Area ISO Survey - South 1 (ELAIS-S1) fields. These methods were trained on both expert cross-identifications of CDFS from the initial ATLAS data release and crowdsourced cross-identifications of CDFS from Radio Galaxy Zoo, and the trained models were used to cross-identify sources in CDFS and ELAIS-S1. The expert-trained methods generally outperform the Radio Galaxy Zoo-trained methods, but the Radio Galaxy Zoo-trained methods still perform reasonably well. Crowdsourced labels thus provide a valuable training set for future work. Cross-identification performance with our method is limited by a lack of radio morphology information as our assumptions require that radio sources are isolated from other sources. This assumption breaks down particularly with high surface brightness sensitivity surveys. We suggest that source identifications and radio morphology information would resolve this problem.

Key words: galaxies: active – galaxies: clusters – radio continuum: galaxies

1 INTRODUCTION

Next generation radio telescopes such as the Australian SKA Pathfinder (ASKAP; Johnston et al. 2007) and Apertif (Verheijen et al. 2008) will conduct increasingly wide, deep, and high-resolution radio surveys, producing large amounts of data. The Evolutionary Map of the Universe survey (EMU; Norris et al. 2011) using ASKAP and the WODAN survey (Röttgering et al. 2011) using Apertif are expected to detect over 100 million radio components between them, compared to the 2.5 million radio components already known (Banfield et al. 2015).

An important part of processing this data is cross-identifying observed radio emission regions with observations of their host galaxy in surveys at other wavelengths. Radio host cross-identification can be a difficult task. Figure 1 illustrates the different radio emission regions that a host galaxy may have. Compact components are unresolved or point-like regions of radio emission while compact radio sources are one compact component usually consistent with the host galaxy location. Extended radio sources consist of resolved single components and/or multiple or resolved or unresolved components. Fanaroff & Riley type I (FRI; Fanaroff

& Riley 1974) and Fanaroff & Riley type II (FRII) are examples of resolved radio sources. Radio emissions from radio-loud active galactic nuclei (AGN) may be complicated structures not clearly related to the host galaxy. These are often composed of multiple, separate components. This type of AGN are expected to dominate approximately 30 per cent of sources detected by EMU (Norris et al. 2011). Small surveys of a few thousand sources such as the Australia Telescope Large Area Survey (ATLAS; Norris et al. 2006; Middelberg et al. 2008) can be cross-identified manually, but this is impractical for larger surveys.

One approach to cross-identification is crowdsourcing, where volunteers cross-identify radio components (both resolved and unresolved; see Figure 1) with their host galaxy. This is the premise of Radio Galaxy Zoo¹ (Banfield et al. 2015), a citizen science project hosted on the highly successful Zooniverse platform (Lintott et al. 2008). Volunteers are shown images of the radio sky, and are asked to identify the radio components in each image that correspond to the same host galaxy. They are then asked to cross-identify

¹ <https://radio.galaxyzoo.org>

the radio source with its host galaxy in a corresponding infrared image. A more thorough explanation of the project can be found in [Banfield et al. \(2015\)](#). The first data release for Radio Galaxy Zoo will provide a large dataset of over 75 000 radio host cross-identifications and radio source morphologies ([Wong et al. prep](#)). While this is a much larger number of visual cross-identifications than have been made by experts (e.g., [Taylor et al. 2007](#); [Gendre & Wall 2008](#); [Grant et al. 2010](#); [Norris et al. 2006](#); [Middelberg et al. 2008](#)) it is still far short of the millions of radio sources expected to be detected in upcoming radio surveys.

Automated algorithms have been developed for this cross-identification problem. [Fan et al. \(2015\)](#) developed a method of cross-identification using Bayesian hypothesis testing, fitting a three-component model to extended radio sources. This was achieved under the assumption that extended radio sources are composed of a core radio component and two lobe components. The core radio component is coincident with the host galaxy, so cross-identification amounts to finding the galaxy coincident with the core radio component in the most likely model fit. This method is easily extended to use other, more complex models, but it is purely geometric. The model does not incorporate other information such as the physical properties of the potential host galaxy. Additionally, there may be new classes of radio sources detected in future surveys like EMU which do not fit the model. [Weston et al. \(prep\)](#) developed a modification of the likelihood ratio method of cross-identification ([Richter 1975](#)) for application to ATLAS and EMU. This method does well on single (resolved or compact) radio sources with a 69 and 71 per cent success rate in the two ATLAS fields, but does not currently handle extended multiple component radio sources ([Norris 2017](#)). A hope is that machine learning techniques can be developed for cross-identification problem. Machine learning describes a class of methods that learn approximations to functions, and so if the cross-identification task can be cast as a machine learning problem, data sets such as that provided by Radio Galaxy Zoo can be generalised to work on data unseen by the original cross-identifiers.

We have developed a machine learning approach for radio host cross-identification. This approach casts the cross-identification task as binary classification, which is very well-understood in machine learning literature and can be solved using a variety of standard, well-known techniques. We have trained these methods on expert cross-identifications and cross-identifications from Radio Galaxy Zoo. In [section 2](#) we describe the data we use to train our methods. In [section 3](#) we discuss how we cast the radio host galaxy cross-identification problem as a machine learning problem. In [section 4](#) we present results of applying our method to ATLAS observations of the *Chandra* Deep Field - South (CDFS) and in [subsection 4.2](#) we apply the cross-identifiers trained on CDFS to the ESO Large Area ISO Survey - South 1 (ELAIS-S1) field. Our data and code are available at <https://radiogalaxyzoo.github.io/atlas-xid>.

2 DATA

We use data from the citizen science project Radio Galaxy Zoo ([Banfield et al. 2015](#)), the Australia Telescope Large Area Survey (ATLAS; [Norris et al. 2006](#); [Franzen et al. 2015](#)), and the *Spitzer* Wide-area Infrared Extragalactic survey (SWIRE; [Lonsdale et al. 2003](#); [Surace et al. 2005](#)).

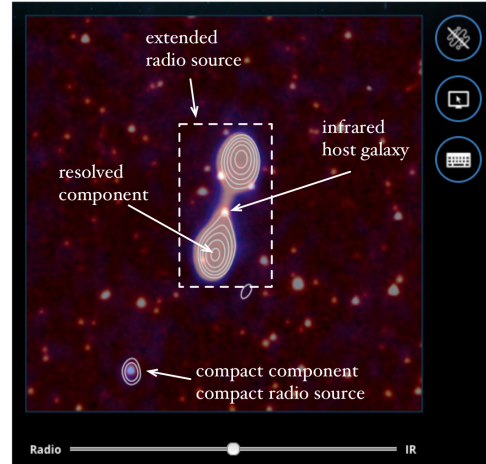


Figure 1. The Radio Galaxy Zoo tutorial image illustrating key definitions used throughout this paper. A colour version of this figure is available online.

Table 1. Catalogues of ATLAS/SWIRE cross-identifications for the CDFS and ELAIS-S1 fields. The method used to generate each catalogue is shown, along with the number of radio components cross-identified in each field.

Catalogue	Method	CDFS	ELAIS-S1
Norris et al. (2006)	Manual	784	0
Middelberg et al. (2008)	Manual	0	1366
Fan et al. (2015)	Bayesian models	784	0
Wong et al. (prep)	Crowdsourcing	2460	0
Weston et al. (prep)	Likelihood ratio	3078	2113

2.1 ATLAS

ATLAS is a pilot survey for the EMU ([Norris et al. 2011](#)) survey, which will cover the entire southern sky and is expected to detect approximately 70 million new radio sources. EMU will be conducted at the same depth and resolution as ATLAS, so methods developed for processing ATLAS data are expected to work for EMU. ATLAS is a wide-area radio survey of the CDFS and ELAIS-S1 fields at 1.4 GHz with a sensitivity of 14 and 17 μ Jy on CDFS and ELAIS-S1 respectively. CDFS covers 3.6 deg² and contains 3034 radio components above 5σ . ELAIS-S1 covers 2.7 deg² and contains 2084 radio components above 5σ ([Franzen et al. 2015](#)). A number of catalogues have been produced cross-identifying radio components in ATLAS with host galaxies in SWIRE. These catalogues are summarised in [Table 1](#).

2.2 SWIRE

SWIRE ([Lonsdale et al. 2003](#); [Surace et al. 2005](#)) is a wide-area infrared survey at the four IRAC wavelengths 3.6, 4.5, 5.8, and 8.0 μ m, with 5σ noise levels of 7.3, 9.7, 27.5, and 32.5 μ Jy respectively ([Lonsdale et al. 2003](#)). It covers eight fields, including CDFS and ELAIS-S1. SWIRE is the source of infrared observations for cross-identification with ATLAS. SWIRE catalogues 221,535 infrared objects in CDFS and 186,059 infrared objects in ELAIS-S1.

2.3 Radio Galaxy Zoo

Radio Galaxy Zoo asks volunteers to cross-identify radio components with their infrared host galaxies. There are a total of 177 461

radio components in Radio Galaxy Zoo, sourced from ATLAS and Faint Images of the Radio Sky at Twenty-Centimeters (FIRST; White et al. 1997). These components are cross-identified to host galaxies detected in SWIRE and by the *Wide-Field Infrared Survey Explorer* (WISE; Wright et al. 2010), respectively. Volunteers are presented with a random radio image centred on a radio component, and a corresponding infrared image centred on the same location. The radio image may contain multiple radio components. For ATLAS the radio and infrared images shown to volunteers are 2×2 arcmin, and for FIRST the images are 3×3 arcmin. Classification of these images is a two-step process. First, volunteers select which radio components are part of the same radio source. Second, volunteers associate each radio source with a host galaxy visible in the infrared image. A more detailed description can be found in Banfield et al. (2015) and a full description of how the final catalogue of cross-identifications is generated can be found in Wong et al. (prep).

In this paper we focus on the Radio Galaxy Zoo cross-identifications of the ATLAS CDFS and SWIRE surveys. The two main reasons are: (1) ATLAS is the pilot study for EMU where automated methods like ours will be used; and (2) ATLAS is composed of two fields, so we can train methods on one field (CDFS) and test these methods on the other field (ELAIS-S1) to ensure that our methods are transferable to different areas of the sky observed by the same telescope.

The ATLAS CDFS radio components that appear in Radio Galaxy Zoo are based on the third data release of ATLAS by Franzen et al. (2015). This release provides a catalogue of radio components (both resolved and unresolved) for each ATLAS field. Each radio component was fit with a two-dimensional Gaussian. Depending on the residual of the fit, more than one Gaussian may be fit to one region of radio emission. Each of these Gaussian fits is listed as a radio component in the ATLAS catalogue. The brightest radio component from the multiple Gaussian fit is called the ‘primary component’. Each primary component found in the ATLAS DR3 component catalogue appears in Radio Galaxy Zoo. Non-primary components may appear within the image of a primary component, but do not have their own entry in Radio Galaxy Zoo. We will henceforth only discuss the primary components.

We use the ATLAS cross-identifications from a preliminary version of the Radio Galaxy Zoo first data release catalogue with no filtering based on volunteer agreement. The final catalogue will introduce a cut on the percentage of volunteers who agreed with a cross-identification, and will hence have more accurate cross-identifications.

3 METHOD

3.1 Cross-identification as binary classification

We focus on the problem of host galaxy cross-identification without reference to radio morphology. Given a radio component, we want to find the corresponding host galaxy as a citizen scientist would using Radio Galaxy Zoo. The input is thus a radio image and an infrared image with a given radius. We choose a radius of 1 arcmin to match the size of the images used by Radio Galaxy Zoo. We make the assumption that each radio image represents a single, complex extended source. This is not the case in general and a radio image may contain many different radio sources with unique host galaxies. If we had access to radio morphology information, we could use this to isolate relevant components and hence more accurately

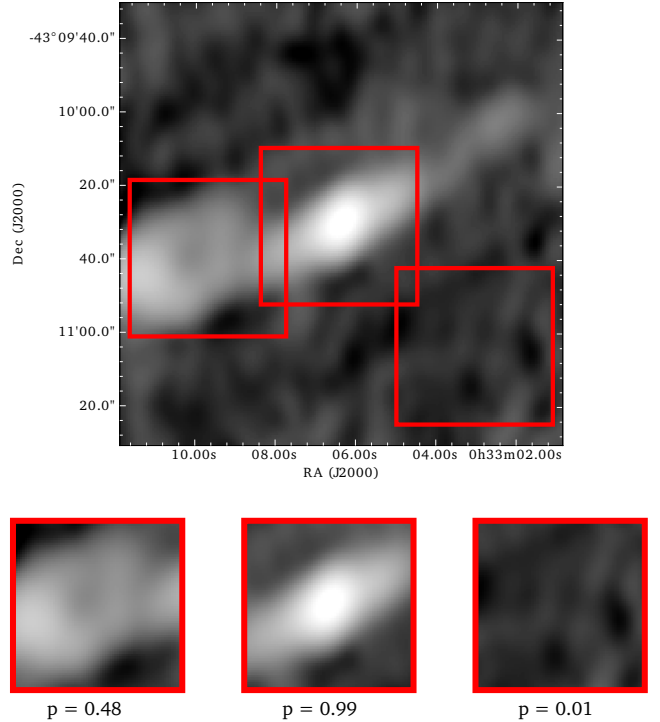


Figure 2. An example of localising the host galaxy of a radio source using our method. This image is from ATLAS and is centred on $\alpha = 00^{\text{h}}33^{\text{m}}06.36^{\text{s}}$, $\delta = -43^{\circ}10'30.1''$. Boxes represent 32×32 pixel windows centred on various locations in the image. The corresponding image patches centred from each window are shown beneath. The image patch of each window represents the location the window is centred on. The probabilities of each patch coinciding with the host galaxy would then be estimated by a classification model. The probabilities shown are for illustration only. In this example, the centre window of the image would be chosen as the location of the host galaxy, as the window centred on it has the highest probability.

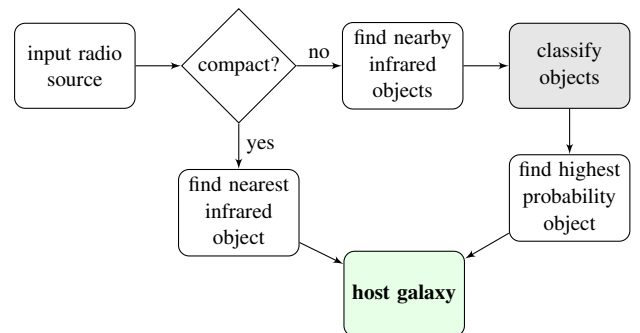


Figure 3. A cross-identification method employing a binary classifier. As input we accept a radio source. If the source is compact, we select the nearest infrared object as the host galaxy. If the source is resolved, we classify all infrared objects nearby within radius R and select the highest probability object as the host galaxy. The grey box is the classifier, which can be any binary classifier that outputs a probability.

cross-identify sources, but this is beyond the scope of this paper. We also note that this assumption will bias our results against radio emission that extends beyond the 1 arcmin cutout size, however, these are uncommon in ATLAS with 9 radio sources identified by Norris et al. (2006) as having components further than 1 arcmin apart (and of these, only 5 are in Radio Galaxy Zoo). The radio cross-identification task then amounts to locating the host galaxy within the associated radio and infrared images. This is formalised as an object localisation problem: given a radio image and an infrared image centred on a radio component, locate the host galaxy of the radio component.

A common approach to object localisation is to estimate the probability that each location in the image is coincident with the desired object. In practice this amounts to taking fixed-size windows of the image centred on each pixel and estimating the probability that each window is centred on the object. The location associated with the highest probability window is then assumed to be the location of the object. Applying this to radio host cross-identification, we consider windows of a radio/infrared image and estimate the probability that each window is centred on the host galaxy. This approach is illustrated in Figure 2. In computer vision a sliding window is commonly used: every possible patch in the image is considered (e.g. Rowley et al. 1996).

This task can be made more efficient if we have a prior on the location of the object we are localising. For our prior, we assume that the host galaxy is always visible in the infrared and thus we only need consider windows centred on infrared sources. This assumption usually holds, except for a rare class of infrared-faint radio sources. Norris et al. (2006) found 22 such radio sources in their sample of 784 bright ATLAS components in CDFS. This leads us to a binary classification task: given an infrared source, compute the probability that it is a host galaxy. To find the host galaxy given a radio source, classify each galaxy within 1 arcmin of the source and select the galaxy with the highest probability of being a host galaxy. This is a good formulation as binary classification is a very common problem in machine learning and there are many different methods readily available to solve it.

Solving the radio cross-identification task amounts to modelling a function f from infrared sources X to the probability that an infrared source belongs to a binary class in $\mathcal{Y} = \{0, 1\}$:

$$f(x) := \Pr(\mathcal{Y} = 1 \mid X = x) \quad . \quad (1)$$

The space of infrared sources X needs to be encoded as a vector for the models we will use. We describe this in subsection 3.3. There are many options for modelling f . In this paper we apply three different models: logistic regression, random forests, and convolutional neural networks.

We can improve upon this cross-identification method by filtering out compact radio sources, which are much easier to cross-identify — the nearest SWIRE object may be identified as the host galaxy, or a more complex method such as likelihood ratios may be applied (see Weston et al. prep). A full cross-identification pipeline making use of this alongside a binary classifier is shown in Figure 3.

3.2 Limitations of our approach

Our assumptions are:

- A radio image represents a single radio source and contains exactly one host galaxy.

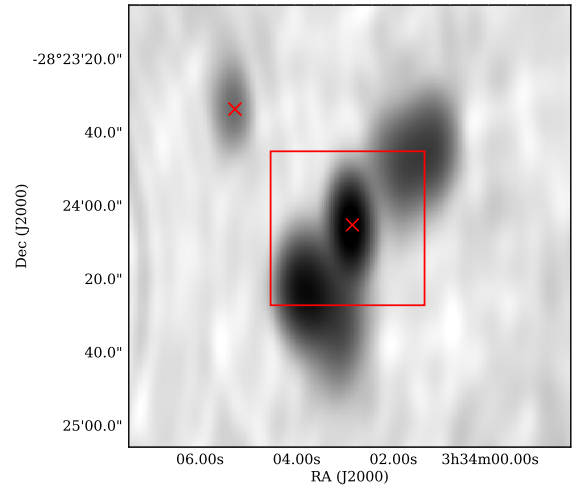


Figure 4. 2'-wide radio image centred on ATLAS3_J033402.87-282405.8C. This radio source breaks our assumption that there are no other radio sources within 1 arcmin of the source. Another radio source is visible to the upper-left. Host galaxies found by Radio Galaxy Zoo volunteers are shown by a cross.

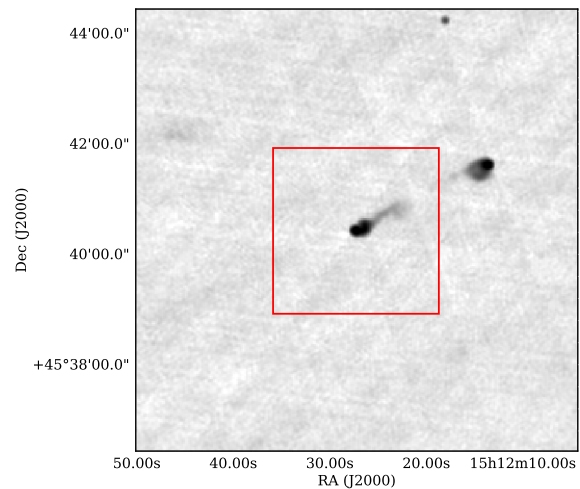


Figure 5. A 8'-wide radio image from FIRST, centred on FIRSTJ151227.2+454026. The 3'-wide red box indicates the boundaries of the image of this radio component shown to volunteers in Radio Galaxy Zoo. This radio source breaks our assumption that the whole radio source is visible in the chosen radius. As one of the lobes of the radio source is outside of the image, a volunteer (or automated algorithm) looking at the 3'-wide image may be unable to determine that this is a radio double or locate the host galaxy.

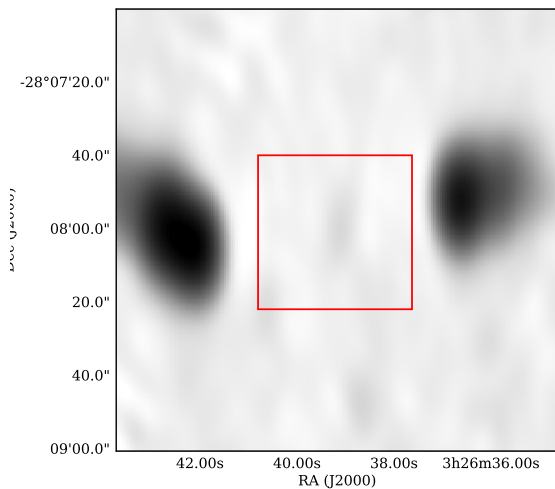


Figure 6. A radio image centred on $\alpha = 03^{\text{h}}26^{\text{m}}39.12^{\text{s}}$, $\delta = -28^{\circ}07'58.80''$. This is an example of a radio source where the window centred on the host galaxy, shown as a rectangle, does not contain enough radio information to correctly identify the galaxy as the host.

- The host galaxy of a radio component is within 1 arcmin of the component.
- The host galaxy appears in the SWIRE catalogue.

The key problem with this approach is our assumption that the radio sky within 1 arcmin radius contains only one, complete radio source. The problem is two-fold: This radius may contain multiple sources, or it may not contain the entirety of the source. If the radius contains multiple sources then there will also be multiple hosts in our input images (which breaks our assumption that there is only one); even a perfect classifier can only accurately cross-identify *one* host in an image with multiple. An example of a radio source that breaks the assumption in this way is shown in Figure 4. If the radius does not contain the whole source, then we are missing radio information useful for finding the host galaxy. This is a difficult problem even for non-automated methods as radio sources can be extremely wide — for example, Radio Galaxy Zoo found a radio giant that spanned over three different images presented to volunteers and the full source was only cross-identified by the efforts of citizen scientists (Banfield et al. 2015). An example of a radio image where part of the radio source is outside the search radius is shown in Figure 5. The problems are in opposition to each other: To reduce the number of sources in an input image, we can reduce the image radius, but this increases the chance that we will miss relevant radio source information, and vice versa.

Our assumptions impose an upper bound on how well we can cross-identify radio sources, which can be estimated by considering how accurately a *perfect* binary classifier cross-identifies radio sources under our method. Using the expert labels (??) to classify SWIRE objects to 100 percent accuracy on the binary classification task results in a cross-identification accuracy of (96.74 ± 1.76) percent across the four CDFS quadrants under our assumptions, which we take as an upper bound on our cross-identification accuracy.

A related issue is that we need to choose a window size for the image representations of each SWIRE object. If this image is too small, radio emission may extend past the edges of the window,

and it may be impossible to identify the galaxy as a host galaxy. If the image is too large, then too much information will be included and it will be difficult or computationally expensive to classify. We chose a window size of 32×32 pixels, shown as a red rectangle in Figure 2 and Figure 6.

3.3 Feature vector representation of infrared sources

Binary classification methods require that the inputs to be classified are represented by an array of real values called *feature vectors*. We thus need to choose a feature vector representation of our candidate host galaxies.

Infrared observations of the CDFS field are taken from SWIRE. We use the CDFS Fall '05 SWIRE catalogue (Surace et al. 2005) to generate candidate hosts to classify. Radio observations of the CDFS field are taken from ATLAS. The radio image is taken to be a 32×32 pixel image from ATLAS, centred on the candidate host location from SWIRE. As the visual appearance of objects in SWIRE are likely irrelevant to whether the galaxy is a host galaxy, the candidate host location encodes almost all information we could gain from the infrared image. We therefore do not use the infrared image for object localisation.

We represent each candidate host as 1034 real-valued features. For a given candidate host, these features are:

- the logarithm of the ratio of fluxes of the candidate host in the four IRAC wavelengths;
- the stellarity index of the host in both 3.6 and 4.5 μm ;
- the flux of the host in 3.6 μm ;
- the radial distance between the candidate host and the nearest radio component in the ATLAS catalogue; and
- a 32×32 pixel image from ATLAS, centred on the candidate host.

The infrared fluxes provide insight into the properties of the host galaxy of the radio emission. The 3.6 and 4.5 μm fluxes trace both galaxies with faint polycyclic aromatic hydrocarbon (PAH) emission and elliptical galaxies dominated by old stellar populations. The 5.8 μm flux selects galaxies where the infrared emission is dominated by non-equilibrium emission of dust grains (PAH destruction by the hard UV spectrum of AGN), while the 8.0 μm flux traces strong PAH emission at low redshift (Sajina et al. 2005). The stellarity index represents how likely the object is to be a star rather than a galaxy.

We use the pixels of each 32×32 radio image as independent features for all classifiers, with the convolutional neural network (Section 3.4.2) automatically extracting features that are relevant. There has been limited research on extracting features from radio images. Proctor (2006) describe hand-selected features for radio doubles in FIRST, and Aniyani & Thorat (2017) and Lukic et al. (prep) make use of deep convolutional neural networks which automatically extract features as part of classification. A more comprehensive search of neural network architectures for processing radio images is beyond the scope of this initial study, and is a good avenue for potential improvement in our pipeline.

3.4 Classifiers

We use three different classifiers as our binary classification model: logistic regression, convolutional neural networks, and random forests.

3.4.1 Logistic Regression

Logistic regression is a binary classification model. It is linear in the feature space and outputs the probability that the input has a positive label. The model is (Bishop 2006):

$$f(\mathbf{x}) = \sigma(\mathbf{w} \cdot \mathbf{x} + b) , \quad (2)$$

where $\mathbf{w} \in \mathbb{R}^D$ is a weights vector, $b \in \mathbb{R}$ is a bias term, $\mathbf{x} \in \mathbb{R}^D$ is the feature representation of a candidate host, and σ is the logistic sigmoid function:

$$\sigma(a) = (1 + \exp(-a))^{-1} . \quad (3)$$

The logistic regression model is fully differentiable, and the weights vector \mathbf{w} can therefore be learned using gradient methods.

3.4.2 Convolutional neural networks

Convolutional neural networks (CNN) are a biologically-inspired prediction model for prediction with image inputs. A number of filters are convolved with an input image to produce output images called *feature maps*. These feature maps can then be convolved again with other filters on subsequent layers, producing a network of convolutions. The whole network is differentiable with respect to the values of the filters and the filters can be learned using gradient methods. The final layer of the network is logistic regression, with the convolved outputs as input features. For more detail, see LeCun et al. (subsection II.A 1998). We use KERAS (Chollet et al. 2015) to implement our CNN.

CNNs have recently produced good results on large image-based datasets in astronomy (e.g. Dieleman et al. 2015; Lukic et al. prep). We employ only a simple CNN model in this paper as a proof of concept that CNNs may be used for class probability prediction on radio images. The model architecture we use is shown in Figure 7.

3.4.3 Random Forests

Random forests are an ensemble of decision trees (Breiman 2001). It considers multiple subsamples of the training set, where each bootstrap subsample is sampled with replacement from the training set. For each subsample a decision tree classifier is constructed. The decision tree is built by repeatedly making axis-parallel splits based on individual features. In a random forest the split decision is taken based on a random subset of features. To classify a new data point, the random forest takes the weighted average of all classifications produced by each decision tree.

3.5 Labels

Converting the Radio Galaxy Zoo and Norris et al. (2006) cross-identification catalogues to binary labels for infrared objects is a non-trivial task. A clear problem is that there is no way to capture radio morphology information in binary classification. As a result, we ignore this problem for this paper. Another problem is that there is no way to indicate *which* radio object an infrared object is associated with, only that it is associated with *some* radio object. We make the naïve assumption that any given radio image contains only one host galaxy as the first step in solving this problem.

We then generate positive labels from a cross-identification catalogue. We decide that if an infrared object is listed in the catalogue, then it is assigned a positive label as a host galaxy.

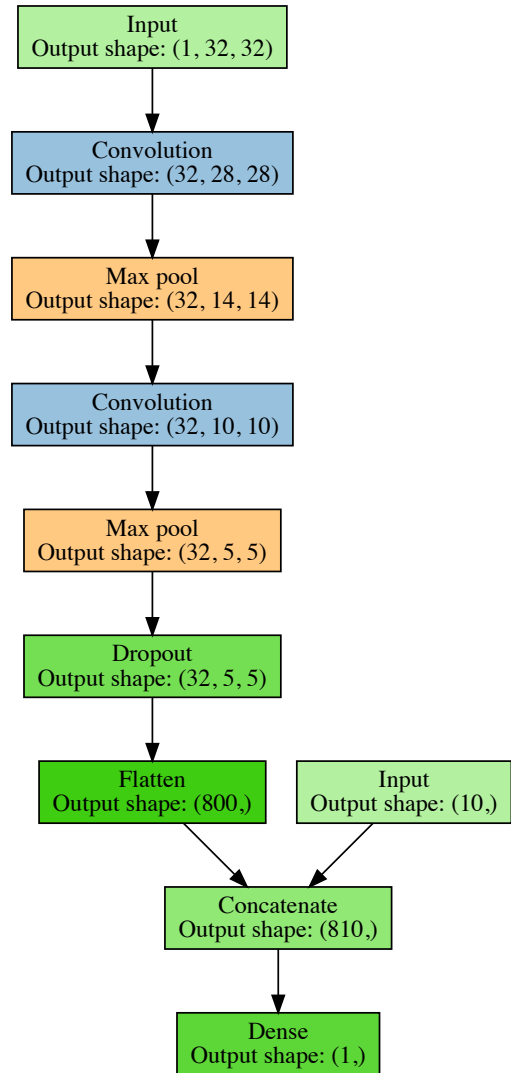


Figure 7. Architecture of our CNN. The merge layer flattens the output of the previous layer and adds the 10 features derived from the candidate host in SWIRE, i.e. the flux ratios, stellarity indices, and distance. The dropout layer drops 25% of its inputs. Diagram based on <https://github.com/dnouri/nolearn>.

In principle we would then assign every other galaxy a negative label. This has some problems — an example is that if the cross-identifier did not observe a radio object (e.g. it was below the signal-to-noise ratio) then the host galaxy of that radio object would receive a negative label. This occurs with Norris et al. (2006) cross-identifications, as these are associated with the first data release of ATLAS. The first data release went to a 5σ flux density level of $S_{1.4} \geq 1 \text{ mJy beam}^{-1}$ (Norris et al. 2006), compared to $S_{1.4} \geq 85 \mu\text{Jy beam}^{-1}$ for the third data release used by Radio Galaxy Zoo (Franzen et al. 2015). The expert labels may therefore disagree with labels from Radio Galaxy Zoo even if they are both plausible. The difference in flux distributions between the expert and Radio Galaxy Zoo cross-identified sources is shown in Fig-

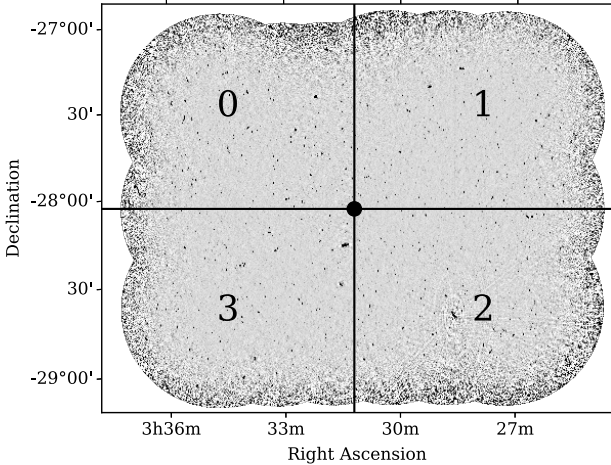


Figure 8. CDFS field training and testing quadrants labelled 0 – 3. The central dot is located at $\alpha = 03^{\text{h}}31^{\text{m}}12^{\text{s}}$, $\delta = -28^{\circ}06'00''$. The quadrants were chosen such that there are similar numbers of radio sources in each quadrant.

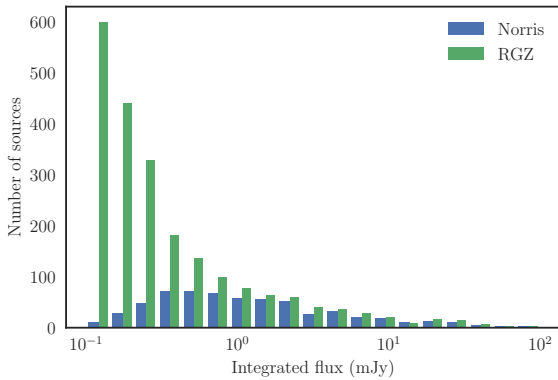


Figure 9. Distribution of integrated flux for radio components cross-identified in ATLAS DR1 (Norris) and by Radio Galaxy Zoo (RGZ). Radio Galaxy Zoo has cross-identified considerably more faint objects.

Figure 9 and the difference in training set size at different flux cutoffs is shown in Figure 10. We train and test our classifiers on infrared objects within a 1 arcmin radius of an ATLAS radio object.

3.6 Experimental Setup

We trained cross-identifiers on radio objects from the ATLAS observations of the CDFS field, using two label sets. One label set was derived from Radio Galaxy Zoo cross-identifications and the other was derived from the Norris et al. (2006) cross-identification catalogue. We refer to these as the Radio Galaxy Zoo labels and the expert labels respectively. We divided the CDFS field into four quadrants for training and testing. The quadrants were centred on $\alpha = 52^{\text{h}}48^{\text{m}}00^{\text{s}}$, $\delta = -28^{\circ}06'00''$ as shown in Figure 8. For each trial, one quadrant was used to draw test examples and the other three quadrants were used for training examples.

We further divided the radio components into compact and resolved. Compact components are cross-identified by fitting a 2D

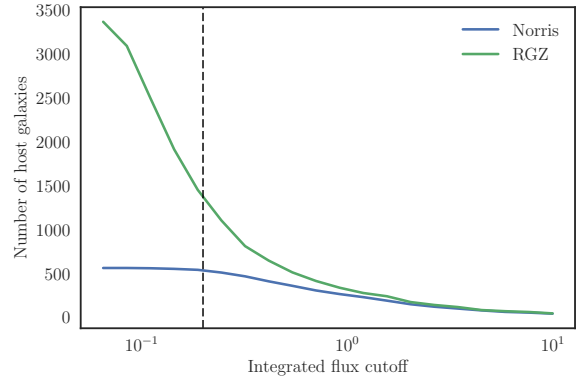


Figure 10. Number of host galaxies in the expert (Norris) and Radio Galaxy Zoo (RGZ) training sets with different integrated radio flux cutoffs. The two training sets converge to the same number of host galaxies as the cutoff is increased.

Gaussian (as in Norris et al. 2006) and we would expect any machine learning approach for host cross-identification to attain high accuracy on this set. Whether a component was resolved was decided based on its flux; a radio component was considered resolved if

$$\ln\left(\frac{S_{\text{int}}}{S_{\text{peak}}}\right) > 2 \sqrt{\left(\frac{\sigma_{S_{\text{int}}}}{S_{\text{int}}}\right)^2 + \left(\frac{\sigma_{S_{\text{peak}}}}{S_{\text{peak}}}\right)^2}, \quad (4)$$

where S_{int} is the integrated flux density and S_{peak} is the peak flux density (Franzen et al. 2015).

Candidate hosts were selected from the SWIRE catalogue. For a given subset of radio components, all SWIRE objects within 1 arcmin of all radio components in the subset were added to the associated SWIRE subset. In the context of the binary classification task, we refer to SWIRE objects within 1 arcmin of a compact radio component as part of the ‘compact set’, and SWIRE objects within 1 arcmin of a resolved radio component as part of the ‘resolved set’.

Classifiers were trained on all available training data in the training quadrants. Test examples were drawn from the remaining quadrant. To reduce bias in the testing data due to the expert labels being generated from a shallower data release of ATLAS, a SWIRE object was only added to the test set if it was within 1 arcmin of a radio object with a cross-identification in both the Norris et al. (2006) catalogue and the Radio Galaxy Zoo catalogue.

Each classifier was trained on the training examples and used to predict labels for the test examples. The predicted labels were compared to the expert labels and the balanced accuracy was computed. We use balanced accuracy as our accuracy measure due to the highly imbalanced classes — in our total set of SWIRE objects within 1 arcmin of an ATLAS object, only 4 percent have positive labels. Only examples within 1 arcmin of ATLAS objects in the first ATLAS data release (Norris et al. 2006) were used to compute accuracy, as these were the only ATLAS objects with expert labels.

We then used the outputs of our classifiers to predict the host galaxy for each radio component cross-identified by both Norris et al. (2006) and Radio Galaxy Zoo. For each SWIRE object within 1 arcmin of the radio component, the probability of the object having a positive label was estimated using the trained binary classifiers. The SWIRE object with the highest probability was chosen

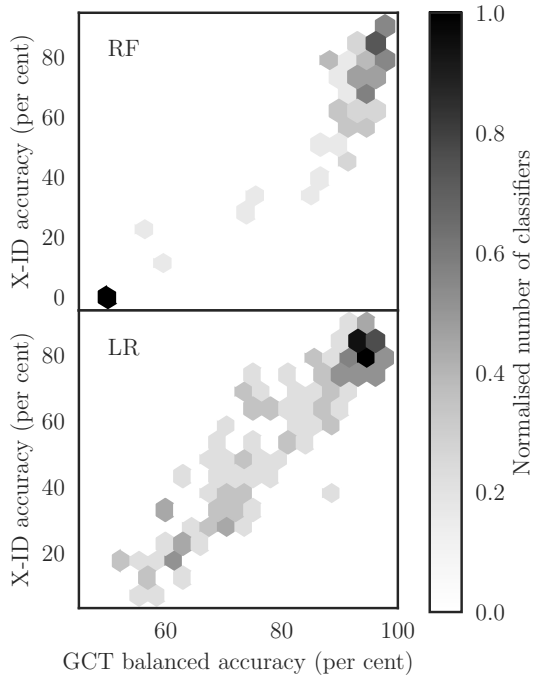


Figure 11. Balanced accuracy on the galaxy classification task (GCT) plotted against accuracy on the cross-identification task (X-ID). RF indicates results from random forests, and LR indicates results from logistic regression. Classifiers were trained on random, small subsets of the training data to artificially restrict their accuracies.

as the host galaxy. The accuracy was then estimated by counting how many predicted host galaxies matched the [Norris et al. \(2006\)](#) cross-identifications.

4 RESULTS

4.1 Application to ATLAS-CDFS

We report two sets of accuracies: first, the balanced accuracies for the task of classifying individual SWIRE objects as host galaxies or not host galaxies (the ‘galaxy classification task’); and second, the accuracies for the task of cross-identifying ATLAS radio components with their host galaxies (the ‘cross-identification task’). Balanced accuracy is a useful metric for evaluating classifiers on classification tasks with highly imbalanced classes like we have here, and we find that the balanced accuracy on the galaxy classification task is correlated with the accuracy on the cross-identification task for logistic regression classifiers. This correlation is shown in [Figure 11](#).

For the galaxy classification task, we report the balanced accuracy of each classification model and training set in [Figure 12](#) for SWIRE objects within 1 arcmin of (1) compact radio components, (2) resolved radio components, and (3) all radio components. These accuracies averaged across the four quadrants are reported in [Table 2](#). We find that the classifiers trained on expert labels generally outperform those trained on Radio Galaxy Zoo labels on almost all classification models, with random forests trained and

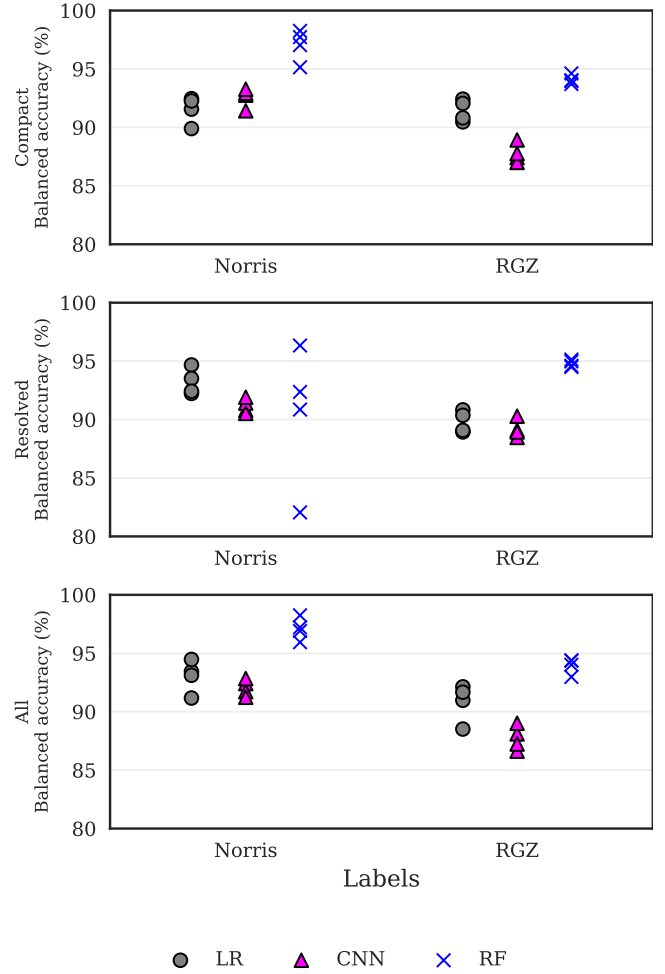


Figure 12. Balanced accuracies for each quadrant in the galaxy classification task. The accuracies of different classifiers are represented by different shapes. The horizontal axis shows which label set was used to train the classifiers: ‘Norris’ represents the expert labels, ‘RGZ’ represents the Radio Galaxy Zoo labels, and ‘RGZ N’ represents the Radio Galaxy Zoo labels applied only to the radio components that appeared in the expert cross-identification catalogue. The three plots show different sets of training and testing data: the ‘compact’ contains only compact radio components, ‘resolved’ only contains resolved radio components, and ‘all’ contains all radio components. The maximum attainable balanced accuracy is 100 percent, when every galaxy label matches the expert labels.

tested on the resolved set as the only exception. This is likely due to the very small size of the expert-labelled resolved set, which had only labelled 98 host galaxies across CDFS compared to 150 host galaxies for the Radio Galaxy Zoo-labelled set. This is supported by the high spread of balanced accuracies for random forests trained on the expert-labelled resolved set, which have a standard deviation of 4.8 percentage points compared to a standard deviation of 0.6 percentage points when trained on the Radio Galaxy Zoo-labelled resolved set. Despite being lower, the accuracies attained by the classifiers trained on Radio Galaxy Zoo labels are still comparable to those trained on expert labels, with most Radio Galaxy Zoo-trained models achieving accuracies within a few percentage points of the expert-trained models.

We report the probabilities predicted by each classifier for each SWIRE object in [Table A1](#). Probabilities reported for a given

Table 2. Average balanced accuracies across quadrants shown in Figure 12 for CDFS. Uncertainties represent standard deviations. The classifiers were trained and tested on compact sources, resolved sources, and all sources separately. ‘Labeller’ indicates what label set the classifiers were trained on. ‘Norris’ is the expert label set, ‘RGZ’ is the Radio Galaxy Zoo label set. ‘LR’ refers to logistic regression, ‘CNN’ refers to our convolutional neural network, and ‘RF’ refers to random forests.

Labeller	Classifier	Mean ‘Compact’ accuracy (per cent)	Mean ‘Resolved’ accuracy (per cent)	Mean ‘All’ accuracy (per cent)
Norris	CNN	92.6 ± 0.7	91.2 ± 0.5	92.1 ± 0.6
	LR	91.6 ± 1.0	93.2 ± 1.0	93.1 ± 1.2
	RF	97.5 ± 0.6	90.9 ± 4.8	96.7 ± 2.3
RGZ	CNN	87.8 ± 0.7	89.2 ± 0.7	87.7 ± 0.9
	LR	91.4 ± 0.8	89.8 ± 0.8	90.8 ± 1.4
	RF	94.0 ± 0.6	94.7 ± 0.6	94.1 ± 0.3

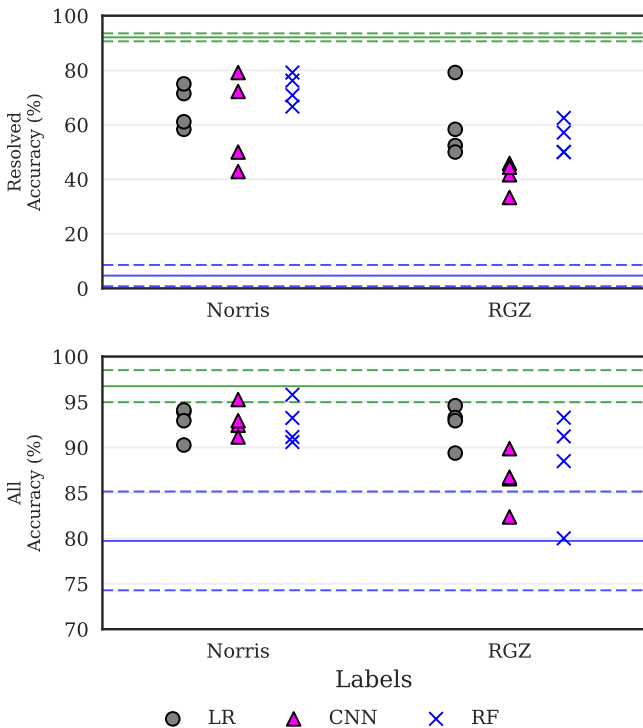


Figure 13. Accuracies for each quadrant in the cross-identification task. Markers and axes are as in Figure 12. The upper horizontal line indicates the accuracy of a ‘perfect’ classifier on the cross-identification task, with the surrounding dashed lines indicating the standard deviation across CDFS quadrants. A ‘perfect’ classifier simply reads the expert labels. This represents the maximum attainable cross-identification accuracy under our assumptions. The lower horizontal line indicates the average accuracy of 15 random classifiers, and hence represents the minimum accuracy we would expect to attain. The surrounding dashed lines indicate the standard deviation across CDFS quadrants and classifiers.

object were predicted by classifiers tested on the quadrant containing that object.

For the cross-identification task, we report the accuracy of each classification model and training label set in Figure 13. We report the averaged cross-identification accuracies across all four quadrants in Table 3. Logistic regression produces the most consistent performance across label sets, with very little variation in accuracy between expert-trained and Radio Galaxy Zoo-trained logistic regression. Random forests perform well when trained on expert labels, but, unlike in the galaxy classification task, perform poorly

Table 3. Average cross-identification accuracies across quadrants. Uncertainties represent standard deviations across quadrants and, for the ‘random’ classifier, 25 random classifiers. Columns and classifiers as defined in Table 2. A ‘perfect’ classifier reads the expert labels directly and hence represents the maximum attainable accuracy on the test set under our assumptions. A ‘random’ classifier generates uniform random class probabilities and hence represents the expected minimum attainable accuracy with our cross-identification method.

Labeller	Classifier	Mean ‘Resolved’ accuracy (per cent)	Mean ‘All’ accuracy (per cent)
Norris	CNN	61.1 ± 15.1	93.0 ± 1.5
	LR	66.5 ± 6.9	92.8 ± 1.5
	RF	67.2 ± 5.9	92.4 ± 1.2
	Perfect	92.1 ± 1.5	96.7 ± 1.8
	Random	4.7 ± 3.9	79.7 ± 5.4
RGZ	CNN	41.3 ± 4.9	86.4 ± 2.7
	LR	60.0 ± 11.5	92.6 ± 1.9
	RF	57.0 ± 6.1	87.1 ± 2.6

when trained on Radio Galaxy Zoo labels. The high spread of accuracies associated with expert-trained random forests in the galaxy classification task is not reflected in the accuracies of the cross-identification task. Convolutional neural networks outperform the other expert-trained methods on the dataset containing all sources, in contrast to their relatively low performance on the galaxy classification task.

The predicted cross-identification for each ATLAS object is reported in Table A2. As with SWIRE predicted probabilities, reported cross-identifications for a given object were generated by classifiers tested on the quadrant containing that object.

In Figure 14 we show examples of where the incorrect cross-identification was selected by the CNN trained on expert labels. We find that radio components with multiple possible SWIRE host galaxies, radio images with multiple radio sources, and radio images with radio components extending beyond the edge of the image are problematic for the classifier. We also note in (c), (d), and (h) that the classifier tends to mistake radio lobes as core emission and thus assigns high probabilities to SWIRE objects near a lobe. This is likely due to the large number of compact objects in the training data — at the resolution of ATLAS, most objects in CDFS appear compact (or very close to compact).

We have noted in subsection 3.5 that the test set of expert labels, derived from the initial ATLAS data release, was less deep than the third data release used by Radio Galaxy Zoo and this paper, introducing a source of label noise in the testing labels. Specifically, true host galaxies may be misidentified as non-host galaxies if the associated radio source was below the 5σ flux density limit in AT-

LAS DR1 but not in ATLAS DR3. This has the effect of reducing the accuracy for Radio Galaxy Zoo-trained classifiers.

4.2 Application to ATLAS-ELAIS

We applied the classifiers trained on CDFS to perform cross-identification on the ELAIS-S1 field. Both CDFS and ELAIS were imaged by the same radio telescope and to similar sensitivities and angular resolution for the ATLAS survey. As ELAIS-S1 has been cross-identified with SWIRE host galaxies by [Middelberg et al. \(2008\)](#), we can use these cross-identifications to derive another set of expert labels, and hence determine how accurate our method is. If our method generalises well across different parts of the sky, then we expect CDFS-trained classifiers applied to cross-identification in ELAIS-S1 to perform comparably to application to CDFS. We report the balanced accuracies for our classifier applied to ELAIS-S1 in [Figure 15](#) and [Table 4](#). [Figure 16](#) shows the accuracies in the cross-identification task on ELAIS-S1 and we list averages accuracies in [Table 5](#).

Cross-identification results from ELAIS-S1 are similar to those for CDFS, showing that classifiers trained on CDFS perform reasonably well on ELAIS-S1. One interesting exception is that random forests trained on expert labels perform well on CDFS but poorly on ELAIS-S1. This is not the case for logistic regression or convolutional neural networks trained on expert labels, nor is it the case for random forests trained on Radio Galaxy Zoo. We hypothesise that this is because the ELAIS-S1 cross-identification catalogue ([Middelberg et al. 2008](#)) labelled fainter radio components than the CDFS cross-identification catalogue ([Norris et al. 2006](#)) due to noise from the very bright source ATCDFS_J032836.53-284156.0 in CDFS. Classifiers trained on CDFS expert labels may thus be biased toward brighter radio components compared to ELAIS-S1. Radio Galaxy Zoo uses the third data release of ATLAS ([Franzen et al. 2015](#)) and so classifiers trained on the Radio Galaxy Zoo labels may be less biased toward brighter sources compared to those trained on the expert labels. To test this hypothesis we tested each classifier against test sets with an integrated flux cutoff. A SWIRE object was only included in the test set for a given cutoff if it was located within $1'$ of a radio component with an integrated flux above the cutoff. The balanced accuracies for each classifier at each cutoff are shown in [Figure 17\(a\)](#) and (b) and the distribution of test set size for each cutoff is shown in [Figure 17\(c\)](#). (c) shows that ELAIS-S1 indeed has more faint objects than CDFS, with the flux cutoff for which the two fields reach the same test set size (~ 0.7 mJy) indicated by the dashed vertical line on each plot. For CDFS, all classifiers perform reasonably well across cutoffs, with performance dropping as the size of the test set becomes small. For ELAIS-S1, logistic regression and convolutional neural networks perform comparably across all flux cutoffs, but random forests do not: while random forests trained on Radio Galaxy Zoo labels perform comparably to other classifiers across all flux cutoffs, random forests trained on expert labels show a considerable drop in performance below the dashed line.

4.3 Implications and Lessons

Our main result is that it is possible to cast radio host galaxy cross-identification as a machine learning task for which standard methods can be applied. These methods can then be trained with a variety of label sets derived from cross-identification catalogues. Future work could even combine multiple catalogues or physical priors to boost performance.

The accuracies of our trained cross-identification methods generally fall below the best possible accuracy attainable using our approach, indicated by the green lines in [Figures 13](#) and [16](#). The balanced accuracies attained by our binary classifiers indicate that there is significant room for improvement in classification. The classification accuracy can be improved by better model selection and more training data, particularly for convolutional neural networks. There is a huge variety of ways to build a convolutional neural network, and we have only investigated one architecture. For an exploration of different convolutional neural network architectures applied to radio astronomy, see [Lukic et al. \(prep\)](#). Convolutional neural networks generally require more training data than other machine learning models and we have only trained our networks on a few hundred sources. We would thus expect performance on the classification task to greatly increase with larger training sets.

Another problem is that of the window size used to select radio features. Increasing window size would increase computational expense, but provide more information to the models. Results are also highly sensitive to how large the window size is compared to the size of the radio galaxy we are trying to cross-identify, with large angular sizes requiring large window sizes to ensure that the features contain all the information needed to localise the host galaxy. An ideal implementation of our method would most likely represent a galaxy using radio images taken at multiple window sizes, but this is considerably more expensive.

Larger training sets, better model selection, and larger window sizes would improve performance, but only so far: we would still be bounded above by the estimated “perfect” classifier accuracy. From this point, the performance can only be improved by improving upon our broken assumptions. We detailed these assumptions in [Figure 3.2](#), and we will discuss here how these could be resolved. Our assumption that we have only one radio object in a given radius is easily broken by radio galaxies with large angular size, as these may overlap other radio objects. Source identification may solve this problem. An example of how source identification could help resolve the issue is to identify all sources, then cross-identify the galaxies in order of angular size, excluding previous host galaxies each time. Source identifications may also provide priors on the true host galaxy location: candidate hosts can be generated based on physical models applied to sources, and then these hosts can be selected from using our model. Our assumption that the host galaxy is contained within the search radius could be improved by dynamically choosing the search radius, perhaps based on the angular extent of the galaxy, or the redshift of candidate hosts. Finally, our assumption that the host galaxy is visible in infrared is technically not needed, as the sliding window approach we have employed will still work even if there are no visible host galaxies — instead of classifying candidate hosts, simply classify each pixel in the radio image. The downside of removing candidate hosts is that we are no longer able to reliably incorporate host galaxy information such as colour and redshift, though this could be resolved by treating pixels as potentially invisible candidate hosts with noisy features.

We observe that Radio Galaxy Zoo-trained methods perform comparably to methods trained on expert labels. This shows that the crowdsourced labels from Radio Galaxy Zoo will indeed provide a valuable source of training data for future machine learning methods in radio astronomy.

[Figure 17](#) reveals interesting behaviour of different classifier models at different flux cutoffs. Logistic regression and convolutional neural networks seem relatively independent of flux, with these models performing well on the fainter ELAIS-S1 sources

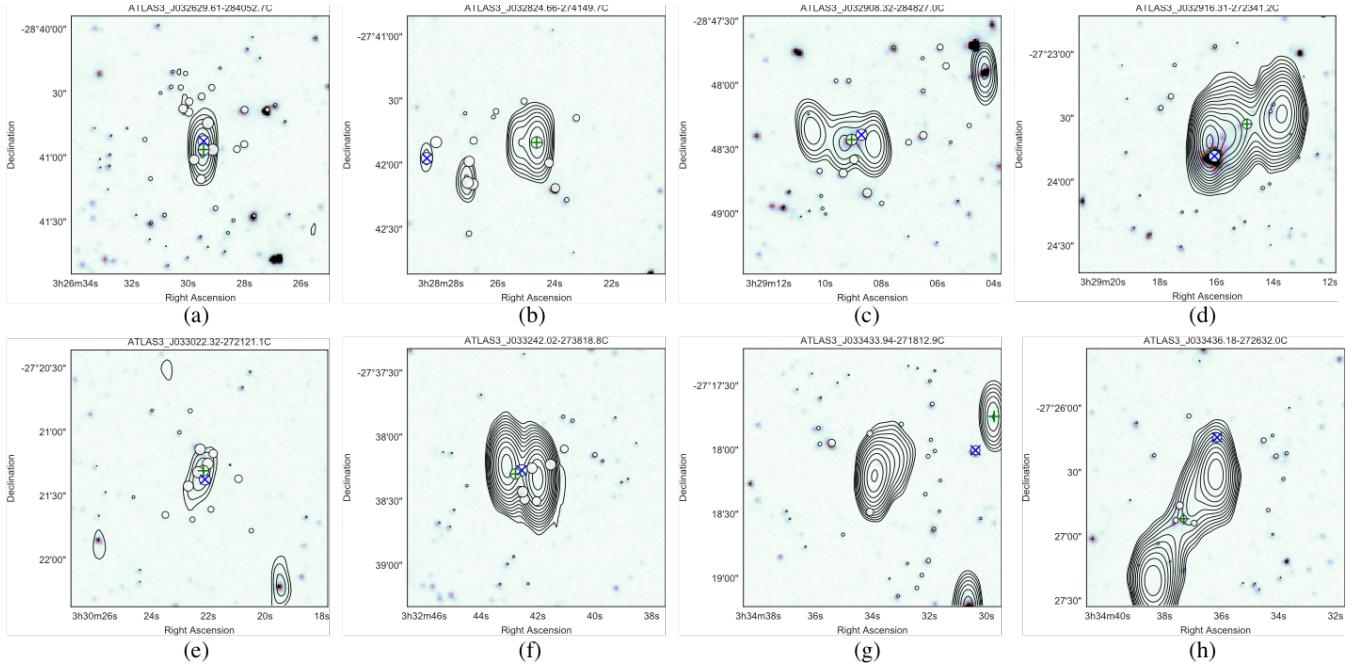


Figure 14. Radio sources in CDFS incorrectly cross-identified by a convolutional neural network trained on expert labels. (a), (e), and (f) have multiple SWIRE host galaxies very close together and the predicted host galaxy is close to the true host galaxy. (b) contains multiple radio sources, and the identified host likely is a host galaxy, but not one labelled in the initial ATLAS data release, as it is a low flux source. In (c), the classifier seems to be distracted by the right lobe of the radio source. This may be due to the large number of compact objects in the training set. (d) and (h) show the classifier failing to correct cross-identify a radio double. This may be due to the window size (both lobes would not be visible in the radio image shown to the classifier) or due to the limited number of radio doubles in the training set. Finally, (g) shows the classifier failing to classify a radio triple. This triple extends outside the 1 arcmin radius used to select candidate hosts.

Table 4. Average balanced accuracies across quadrants shown in Figure 15 for ELAIS-S1. Uncertainties represent standard deviations. The classifiers were applied to compact sources, resolved sources, and all sources separately.

Labeller	Classifier	Mean ‘Compact’ accuracy (per cent)	Mean ‘Resolved’ accuracy (per cent)	Mean ‘All’ accuracy (per cent)
Norris	CNN	94.8 ± 0.2	92.8 ± 0.5	94.4 ± 0.2
	LR	94.6 ± 0.4	93.3 ± 2.0	95.3 ± 0.1
	RF	87.4 ± 5.6	71.5 ± 2.3	86.0 ± 2.6
RGZ	CNN	88.8 ± 0.6	89.9 ± 0.3	88.5 ± 0.3
	LR	93.2 ± 0.5	91.4 ± 0.3	91.6 ± 0.6
	RF	94.9 ± 0.2	95.2 ± 0.3	95.1 ± 0.2

even when they were trained on the generally brighter objects in CDFS. Conversely, random forests were sensitive to the changes in flux distribution between datasets. This shows that not all models behave similarly on radio data, and it is therefore important to investigate multiple models when developing machine learning methods for radio astronomy.

Our methods can be easily incorporated into other cross-identification methods or used as an extra data source for source identification. This is because we have essentially produced a scoring function, which rates galaxies based on the probability that they are a host galaxy. For example, our method could be used to disambiguate between candidate host galaxies selected by model-based algorithms, or used to weight candidate host galaxies while a source identifier attempts to associate radio components. Our method can also be extended using other data sources; for example, information from source identification algorithms could be incorporated into the feature set of candidate host galaxies.

5 SUMMARY

We present a machine learning approach for cross-identification of radio components with their corresponding infrared host galaxy. We trained our methods on expert and crowdsourced cross-identification catalogues, and applied these methods on the CDFS and ELAIS-S1 fields of ATLAS. Expert-trained models generally outperformed the crowdsourced-trained models, but both sets of models performed comparably on the cross-identification task. While our cross-identification performance is not as high as would be ideal, we make no assumptions on the binary classification model used in our methods and so we expect the performance to be improved by further experimentation and model selection. Our method provides a useful framework for generalising cross-identification catalogues to unseen areas of the sky and can be incorporated into existing methods.

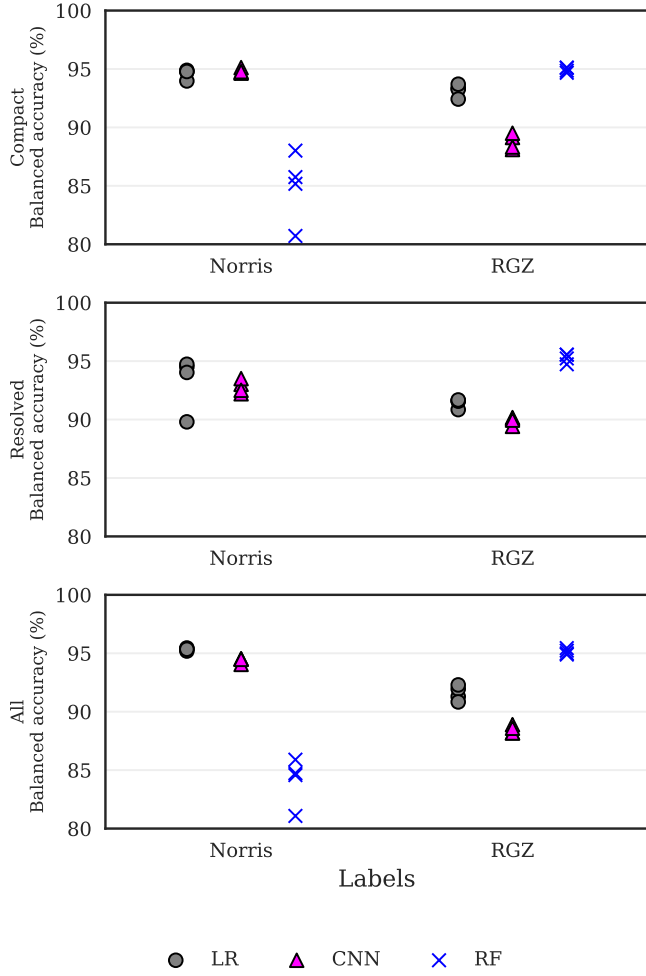


Figure 15. Balanced accuracies for each quadrant in the galaxy classification task on the ELAIS-S1 field. Markers and axes are as in Figure 12. The classifiers in Figure 12 were applied to the whole ELAIS-S1 field, and performance was measured against ground truth labels from Middelberg et al. (2008). Accuracies for random forests trained on the expert-labelled resolved set are not shown as they were considerably lower than all other accuracies.

Table 5. Average cross-identification accuracies for ELAIS-S1. Uncertainties represent standard deviations across classifiers trained on CDFS quadrants. Columns and classifiers as defined in Table 2. A ‘perfect’ classifier reads the expert labels directly and hence represents the maximum attainable accuracy under our assumptions. A ‘random’ classifier generates uniform random class probabilities and hence represents the expected minimum attainable accuracy with our cross-identification method.

Labeller	Classifier	Mean ‘Resolved’ accuracy (per cent)	Mean ‘All’ accuracy (per cent)
Middelberg	Perfect	85.8 ± 0.6	92.4 ± 0.3
	Random	4.7 ± 1.0	50.7 ± 0.5
Norris	CNN	68.1 ± 2.5	85.6 ± 0.7
	LR	79.6 ± 0.9	90.2 ± 0.6
	RF	76.1 ± 3.4	90.4 ± 1.5
RGZ	CNN	59.1 ± 3.2	76.5 ± 0.5
	LR	75.8 ± 2.1	88.3 ± 0.8
	RF	75.2 ± 2.9	86.2 ± 1.3

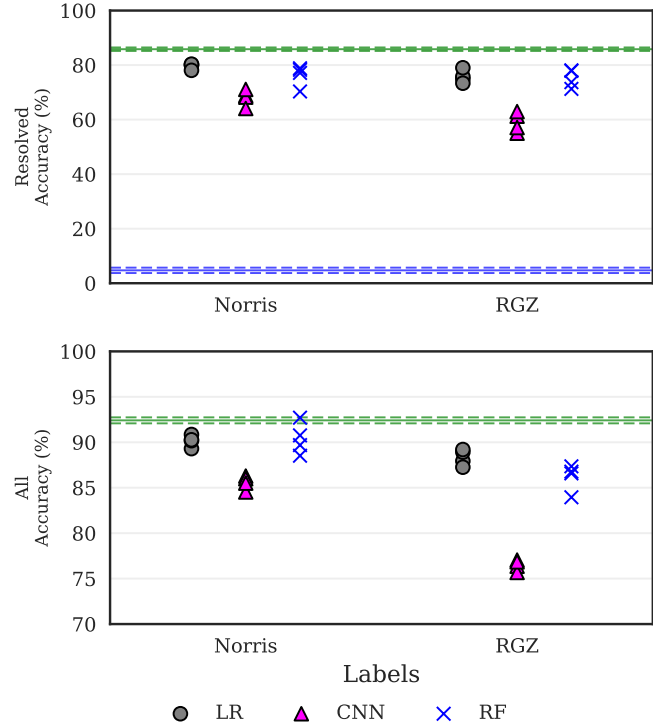


Figure 16. Accuracies for each quadrant in the cross-identification task on the ELAIS-S1 field. Markers and axes are as in Figure 12. The classifiers in Figure 13 were applied to the whole ELAIS-S1 field, and performance was measured against ground truth labels from Middelberg et al. (2008).

References

- Aniyan A. K., Thorat K., 2017, *ApJS*, **230**, 20
Banfield J. K., et al., 2015, *MNRAS*, **453**, 2326
Bishop C. M., 2006, Pattern recognition and machine learning. Springer
Breiman L., 2001, Machine Learning, 45, 5
Chollet F., et al., 2015, Keras, <https://github.com/fchollet/keras>
Dieleman S., Willett K. W., Dambre J., 2015, *MNRAS*, **450**, 1441
Fan D., Budavari T., Norris R. P., Hopkins A. M., 2015, *MNRAS*, **451**, 1299
Fanaroff B. L., Riley J. M., 1974, *MNRAS*, **167**, 31P
Franzen T. M. O., et al., 2015, *MNRAS*, **453**, 4020
Gendre M. A., Wall J. V., 2008, *MNRAS*, **390**, 819
Grant J. K., Taylor A. R., Stil J. M., Landecker T. L., Kothes R., Ransom R. R., Scott D., 2010, *ApJ*, **714**, 1689
Johnston S., et al., 2007, *PASA*, **24**, 174
LeCun Y., Bottou L., Bengio Y., Haffner P., 1998, Proceedings of the IEEE, 86, 2278
Lintott C. J., et al., 2008, *MNRAS*, **389**, 1179
Lonsdale C. J., et al., 2003, *PASP*, **115**, 897
Lukic V., Banfield J. K., Wong O. I., Brüggem M., in prep., Radio Galaxy Zoo: Compact and extended radio source classification with deep learning
Middelberg E., et al., 2008, *AJ*, **135**, 1276
Norris R. P., 2017, *PASA*, **34**, e007
Norris R. P., et al., 2006, *AJ*, **132**, 2409
Norris R. P., et al., 2011, *PASA*, **28**, 215
Proctor D. D., 2006, *ApJS*, **165**, 95
Richter G. A., 1975, *Astronomische Nachrichten*, **296**, 65
Röttgering H., et al., 2011, *Journal of Astrophysics and Astronomy*, **32**, 557
Rowley H. A., Baluja S., Kanade T., 1996, in Advances in Neural Information Processing Systems. pp 875–881
Sajina A., Lacy M., Scott D., 2005, *ApJ*, **621**, 256
Surace J., et al., 2005, Spitzer Science Centre, California Institute of Tech-

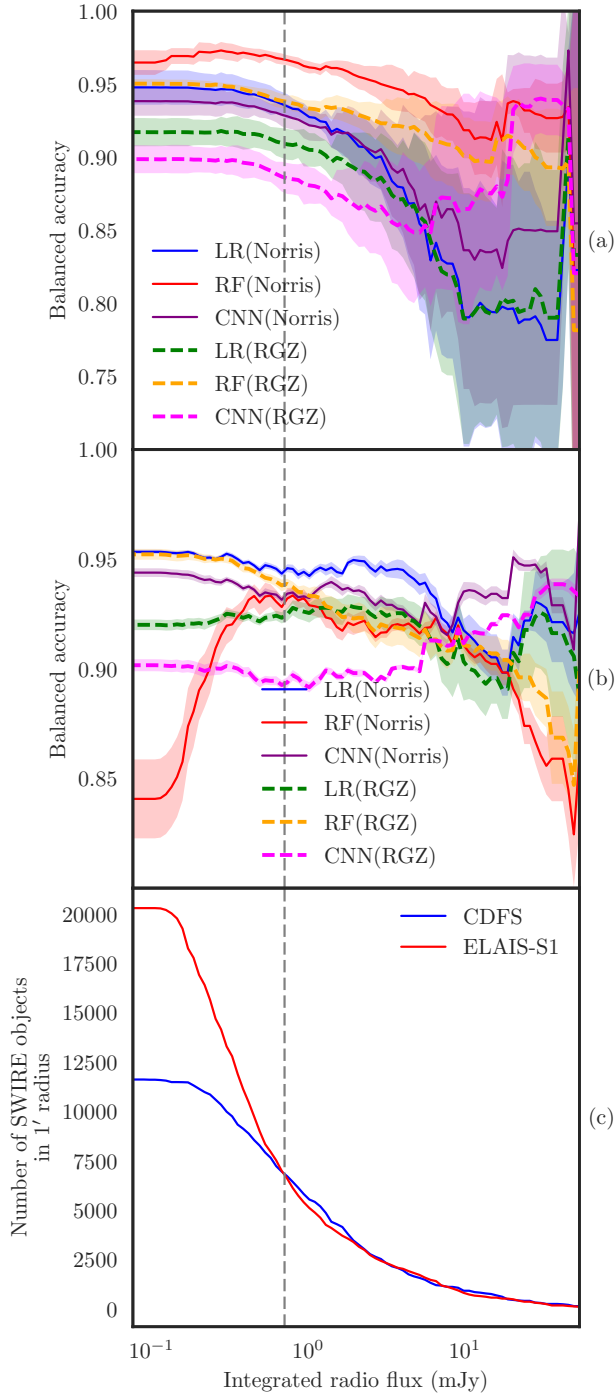


Figure 17. (a) Balanced accuracies of classifiers trained and tested on CDFS with different integrated flux cutoffs for the test set. A SWIRE object is included in the test set if it is within $1'$ of a radio component with greater integrated flux than the cutoff. Different coloured lines indicate different classifier/training labels combinations, where LR is logistic regression, RF is random forests, CNN is convolutional neural networks, and Norris and RGZ are the expert and Radio Galaxy Zoo label sets respectively. Filled areas represent standard deviations across CDFS quadrants. (b) Balanced accuracies of classifiers trained on CDFS and tested on ELAIS-S1. (c) A cumulative distribution plot of SWIRE objects associated with a radio object with more integrated flux than the cutoff.

nology, Pasadena, CA
 Taylor A. R., et al., 2007, *ApJ*, **666**, 201
 Verheijen M. A. W., Oosterloo T. A., van Cappellen W. A., Bakker L., Ivashina M. V., van der Hulst J. M., 2008, in AIP Conf. Ser. 1035, The Evolution of Galaxies Through the Neutral Hydrogen Window. pp 265–271
 Weston S. D., Seymour N., Gulyaev S., Norris R. P., Banfield J., Vaccari M., Hopkins A. M., Franzen T. M. O., in prep., Automated Cross-matching Radio to Infrared Surveys with the LRPY algorithm
 White R. L., Becker R. H., Helfand D. J., Gregg M. D., 1997, *ApJ*, **475**, 479
 Wong O. I., et al., in prep., Radio Galaxy Zoo Data Release 1: morphological classifications of 100,000 FIRST radio sources
 Wright E. L., et al., 2010, *AJ*, **140**, 1868

APPENDIX A: CATALOGUES

The predicted probabilities for each SWIRE object are listed in [Table A1](#). A full version of the table is available electronically. The columns are defined as follows:

- *Column 1* – SWIRE. The SWIRE object name given in [Surace et al. \(2005\)](#).
- *Column 2* – RA. The right ascension of the SWIRE object in decimal degrees (J2000).
- *Column 3* – Dec. The declination of the SWIRE object in decimal degrees (J2000).
- *Columns 4 – 21* – $C(A / B)$. The predicted probability that the SWIRE object is a host galaxy according to classifier C trained on subset B of label set A . ‘LR’ indicates logistic regression, ‘RF’ indicates random forests, ‘CNN’ indicates convolutional neural networks, ‘Norris’ indicates the expert labels, and ‘RGZ’ indicates the Radio Galaxy Zoo labels.

The predicted cross-identifications for each ATLAS object are listed in [Table A2](#). A full version of the table is available electronically. The columns are defined as follows:

- *Column 1* – ATLAS. The ATLAS object name given in [Franzen et al. \(2015\)](#).
- *Column 2* – RA. The right ascension of the ATLAS object in decimal degrees (J2000).
- *Column 3* – Dec. The declination of the ATLAS object in decimal degrees (J2000).
- *Column 4* – Expert. The expert SWIRE cross-identification. This is the SWIRE object associated with the ATLAS object according to either [Norris et al. \(2006\)](#) or [Middelberg et al. \(2008\)](#) for CDFS and ELAIS-S1, respectively.
- *Column 5* – RGZ. The Radio Galaxy Zoo cross-identification. This is the SWIRE object associated with the ATLAS object according to Radio Galaxy Zoo.
- *Column 5* – RGZ radio consensus. The percentage agreement on the radio source according to Radio Galaxy Zoo. Lower numbers tend to indicate more complex sources. For more information on how consensus is generated, see [Wong et al. \(prep\)](#).
- *Column 6* – RGZ IR consensus. The percentage agreement on the host galaxy cross-identification according to Radio Galaxy Zoo. Lower numbers tend to indicate sources that are harder to cross-identify. For more information on how consensus is generated, see [Wong et al. \(prep\)](#).
- *Columns 7 – 25* – $C(A / B)$. The predicted cross-identification according to classifier C trained on subset B of label set A . ‘LR’ indicates logistic regression, ‘RF’ indicates random forests, ‘CNN’ indicates convolutional neural networks, ‘Norris’ indicates the expert labels, and ‘RGZ’ indicates the Radio Galaxy Zoo labels.

This paper has been typeset from a \LaTeX file prepared by the author.

Table A1. Predicted probabilities that each object in SWIRE is a host galaxy. Probabilities are reported for each predictor. $C(A / B)$ indicates the predictor using classifier model C , trained on label set A on data set B . Ellipsis indicates columns that have been omitted. The omitted columns are all other combinations of classifier, label set, and data set. If a SWIRE object does not appear in the table, then it was further than 1 arcmin from an ATLAS object and hence has a predicted probability of zero by our assumptions. Full table electronic.

SWIRE	RA	Dec	CNN(Norris / All)	RF(RGZ / Resolved)
SWIRE3_J032559.15-284724.2	51.4965	-28.7901	0.0001	... 0.1170
SWIRE3_J032559.91-284728.9	51.4996	-28.7914	0.0004	... 0.0000
SWIRE3_J032600.02-284736.9	51.5001	-28.7936	0.0002	... 0.0451
SWIRE3_J032600.13-284637.5	51.5005	-28.7771	0.0004	... 0.0637
SWIRE3_J032600.13-284715.7	51.5006	-28.7877	0.0014	... 0.0153
SWIRE3_J032600.98-284705.4	51.5041	-28.7848	0.0205	... 0.0000
SWIRE3_J032601.03-284711.6	51.5043	-28.7866	0.0699	... 0.0000
SWIRE3_J032601.56-284131.0	51.5065	-28.692	0.0001	... 0.0000
SWIRE3_J032601.60-284207.5	51.5067	-28.7021	0.0000	... 0.0000

Table A2. Predicted host galaxy cross-identifications for each object in ATLAS. Cross-identifications are reported for each predictor, with the predictor listed as in Table A1. The cross-identification given by Norris et al. (2006) for CDFS or Middelberg et al. (2008) for ELAIS-S1 is included in the ‘Expert’ column. For CDFS, the cross-identification given by the Radio Galaxy Zoo consensus is included in the ‘RGZ’ column, along with the corresponding consensus levels for the radio morphology and host cross-identification tasks (see Wong et al. prep, for details on how consensus is calculated). Low radio consensus indicates that the component has multiple nearby components (and thus is more impacted by our assumption that the source is isolated), and low infrared consensus indicates that the host galaxy is unclear in the SWIRE image (possibly due to multiple nearby candidate host galaxies). Full table electronic.

ATLAS	RA	Dec	Expert	RGZ
ATLAS3_J032602.82-284708.1C	51.511734	-28.785575	SWIRE3_J032603.15-284708.5	-
ATLAS3_J032615.49-284629.4C	51.564555	-28.774847	SWIRE3_J032615.41-284630.7	SWIRE3_J032615.41-284630.7
ATLAS3_J032615.55-280559.8C	51.564799	-28.099955	SWIRE3_J032615.52-280559.8	SWIRE3_J032615.52-280559.8
ATLAS3_J032617.35-280710.2C	51.572279	-28.119491	SWIRE3_J032617.89-280707.2	SWIRE3_J032617.89-280707.2
ATLAS3_J032625.13-280909.8C	51.604711	-28.152731	SWIRE3_J032625.19-280910.1	SWIRE3_J032625.19-280910.1
ATLAS3_J032629.10-280650.1C	51.621251	-28.113924	SWIRE3_J032629.13-280650.7	SWIRE3_J032626.74-280636.7
ATLAS3_J032629.61-284052.7C	51.623385	-28.681315	SWIRE3_J032629.54-284055.8	SWIRE3_J032629.54-284055.8
ATLAS3_J032629.92-284753.5C	51.624653	-28.798195	SWIRE3_J032629.81-284754.4	SWIRE3_J032629.81-284754.4
ATLAS3_J032630.66-283657.3C	51.62777	-28.615917	SWIRE3_J032630.64-283658.0	SWIRE3_J032628.56-283744.8
RGZ radio consensus	RGZ IR consensus	CNN(Norris / All)		RF(RGZ / Resolved)
0.4516	0.3214	SWIRE3_J032602.36-284711.5	...	SWIRE3_J032603.60-284627.4
0.2941	0.8000	SWIRE3_J032615.41-284630.7	...	SWIRE3_J032615.41-284630.7
0.5625	0.8333	SWIRE3_J032615.52-280559.8	...	SWIRE3_J032615.52-280559.8
0.4146	1.0000	SWIRE3_J032617.89-280707.2	...	SWIRE3_J032617.89-280707.2
0.3158	0.6667	SWIRE3_J032625.19-280910.1	...	SWIRE3_J032625.19-280910.1
0.3333	1.0000	SWIRE3_J032629.13-280650.7	...	SWIRE3_J032629.13-280650.7
0.2676	1.0000	SWIRE3_J032629.54-284051.9	...	SWIRE3_J032629.54-284051.9
1.0000	0.8571	SWIRE3_J032629.81-284754.4	...	SWIRE3_J032629.81-284754.4
0.3611	0.7308	SWIRE3_J032630.64-283658.0	...	SWIRE3_J032628.56-283744.8